

✓ Importing libraries

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

✓ Mount google drive to fetch the dataset

```
from google.colab import drive
drive.mount('/content/drive')
```

🔗 Mounted at /content/drive

✓ Read dataset

```
df = pd.read_csv('/content/drive/MyDrive/ProjectStage/Liver/Dataset/All_Combined.csv')
```

```
df.sample(5)
```

🔗

	Age	Gender	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Aminotransferase	Aspartate_Am
348	45.0	Male	2.400000	1.100000	168		33
1994	53.0	m	4.664552	3.056918	744.1800508	70.07613446	
944	NaN	NaN	3.210000	0.860000	278.44		61.55

979	NaN	NaN	0.320000	0.140000	91.18	26.76
1648	27.0	m	3.366542	2.044608	469.9292357	260.6680803

✓ Features in dataset

```
df.shape
```

```
↳ (2391, 11)
```

```
df.columns
```

```
↳ Index(['Age', 'Gender', 'Total_Bilirubin', 'Direct_Bilirubin',
        'Alkaline_Phosphotase', 'Alamine_Aminotransferase',
        'Aspartate_Aminotransferase', 'Total_Protiens', 'Albumin',
        'Albumin_and_Globulin_Ratio', 'Dataset'],
        dtype='object')
```

✓ Finding null values

```
# function to convert n and N to null values
```

```
def convert_to_nan(df):
```

```
    columns = df.columns
```

```
    for col in columns:
```

```
        df[col] = df[col].replace({'n': np.nan, 'N': np.nan})
```

```
    return df
```

```
df = convert_to_nan(df)
```

```
df.isna().sum()
```

	0
Age	500
Gender	500
Total_Bilirubin	0
Direct_Bilirubin	0
Alkaline_Phosphotase	22
Alamine_Aminotransferase	5
Aspartate_Aminotransferase	16
Total_Protiens	0
Albumin	0
Albumin_and_Globulin_Ratio	4
Dataset	0

dtype: int64

✓ Removing null values

```
df.drop('Gender', axis=1, inplace=True)
df.drop('Age', axis=1, inplace=True)
```

```
df.isna().sum()
```

	0
Total_Bilirubin	0
Direct_Bilirubin	0
Alkaline_Phosphotase	22
Alamine_Aminotransferase	5
Aspartate_Aminotransferase	16
Total_Protiens	0
Albumin	0
Albumin_and_Globulin_Ratio	4
Dataset	0

dtype: int64

```
# Convert 'Alkaline_Phosphotase' column to numeric, coercing errors to NaN
df['Alkaline_Phosphotase'] = pd.to_numeric(df['Alkaline_Phosphotase'], errors='coerce')
df['Alamine_Aminotransferase'] = pd.to_numeric(df['Alamine_Aminotransferase'], errors='coerce')
df['Aspartate_Aminotransferase'] = pd.to_numeric(df['Aspartate_Aminotransferase'], errors='coerce')
df['Albumin_and_Globulin_Ratio'] = pd.to_numeric(df['Albumin_and_Globulin_Ratio'], errors='coerce')
df['Total_Protiens'] = pd.to_numeric(df['Total_Protiens'], errors='coerce')
df['Albumin'] = pd.to_numeric(df['Albumin'], errors='coerce')

# Now fill NaN values with the mean

df['Alkaline_Phosphotase'] = df['Alkaline_Phosphotase'].fillna(np.mean(df['Alkaline_Phosphotase']))
df['Alamine_Aminotransferase'] = df['Alamine_Aminotransferase'].fillna(np.mean(df['Alamine_Aminotransferase']))
df['Aspartate_Aminotransferase'] = df['Aspartate_Aminotransferase'].fillna(np.mean(df['Aspartate_Aminotransferase']))
df['Albumin_and_Globulin_Ratio'] = df['Albumin_and_Globulin_Ratio'].fillna(np.mean(df['Albumin_and_Globulin_Ratio']))
```

```
df['Total_Protiens'] = df['Total_Protiens'].fillna(np.mean(df['Total_Protiens']))
df['Albumin'] = df['Albumin'].fillna(np.mean(df['Albumin']))
```

```
# df.dropna(inplace=True)
```

```
df.dtypes
```

	0
Total_Bilirubin	float64
Direct_Bilirubin	float64
Alkaline_Phosphotase	float64
Alamine_Aminotransferase	float64
Aspartate_Aminotransferase	float64
Total_Protiens	float64
Albumin	float64
Albumin_and_Globulin_Ratio	float64
Dataset	int64

```
dtype: object
```

```
df.isna().sum()
```

	0
Total_Bilirubin	0
Direct_Bilirubin	0
Alkaline_Phosphotase	0

Alamine_Aminotransferase	0
Aspartate_Aminotransferase	0
Total_Protiens	0
Albumin	0
Albumin_and_Globulin_Ratio	0
Dataset	0

dtype: int64

df.shape

(2391, 9)

✦ Dividing independent and dependent data in X and Y respectively

df.columns

```
Index(['Total_Bilirubin', 'Direct_Bilirubin', 'Alkaline_Phosphotase',
      'Alamine_Aminotransferase', 'Aspartate_Aminotransferase',
      'Total_Protiens', 'Albumin', 'Albumin_and_Globulin_Ratio', 'Dataset'],
      dtype='object')
```

```
X = df.iloc[:, :-1]
X.sample(5)
```

	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Aminotransferase	Aspartate_Aminotransferas
1774	7.973344	3.036862	546.319440	181.562733	348.85248
342	2.600000	1.200000	410.000000	59.000000	57.00000

352	1.000000	0.300000	208.000000	17.000000	15.00000
1510	2.830050	2.859360	592.660559	216.284220	351.17047
1541	2.655904	1.388051	660.737060	136.504003	359.54097

X.columns

```
Index(['Total_Bilirubin', 'Direct_Bilirubin', 'Alkaline_Phosphotase',
      'Alamine_Aminotransferase', 'Aspartate_Aminotransferase',
      'Total_Protiens', 'Albumin', 'Albumin_and_Globulin_Ratio'],
      dtype='object')
```

```
y = df.iloc[:, -1]
y.sample(6)
```

	Dataset
37	0
2170	1
107	0
572	2
480	1
2147	1

dtype: int64

✓ Unique values in target variable (dependent variable)

```
y.unique()
```

```
array([0, 1, 2])
```

```
y.isna().sum()
```

```
np.int64(0)
```

```
y = y.fillna(0)
```

```
y = y.astype(int)
```

```
y.dtype
```

```
dtype('int64')
```

```
y[:5]
```

Dataset	
0	0
1	0
2	0
3	0
4	0

dtype: int64

```
y=y.map({1:1, 2:0, 0:0})
```


y

Dataset	
0	0
1	0
2	0
3	0
4	0
...	...
2386	1
2387	1
2388	1
2389	1
2390	1

2391 rows × 1 columns

dtype: int64

y.value_counts()

count	
Dataset	
1	1668
0	723

dtype: int64

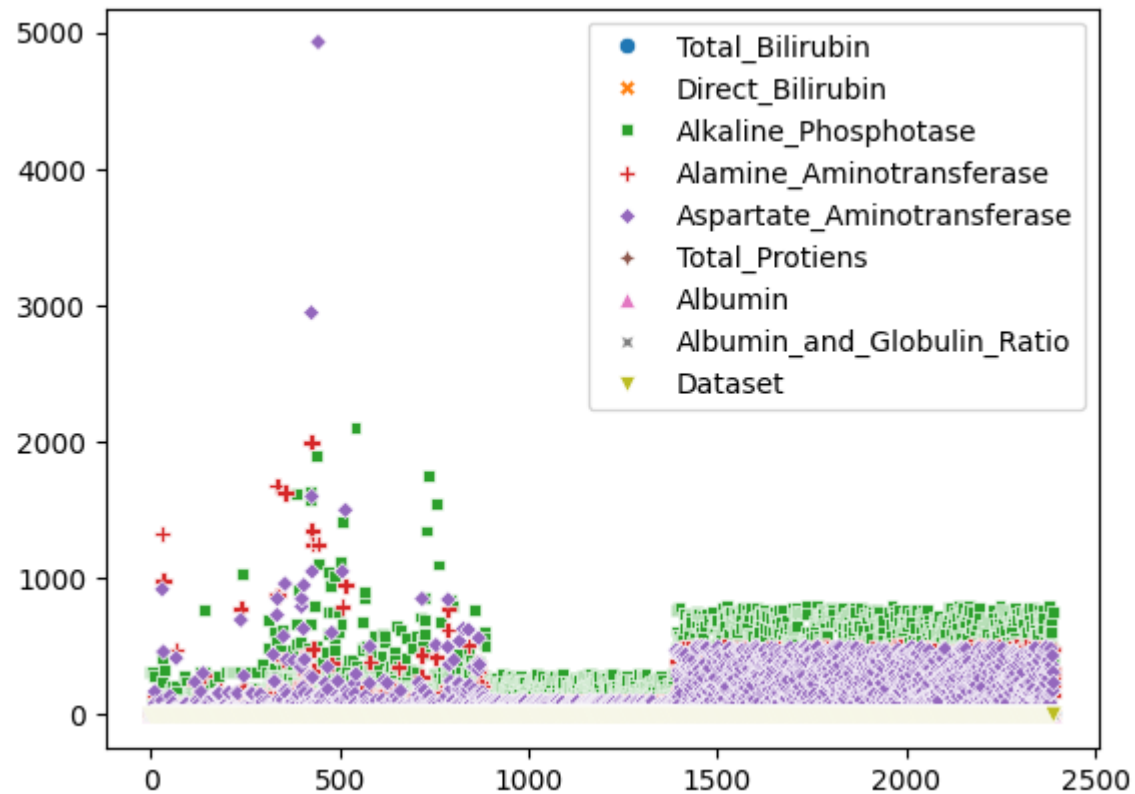
--y, y_counts]

```
y.value_counts()[0], y.value_counts()[1]  
  
(np.int64(723), np.int64(1668))
```

✓ Finding outlier using scatter plot

```
sns.scatterplot(data=df)
```

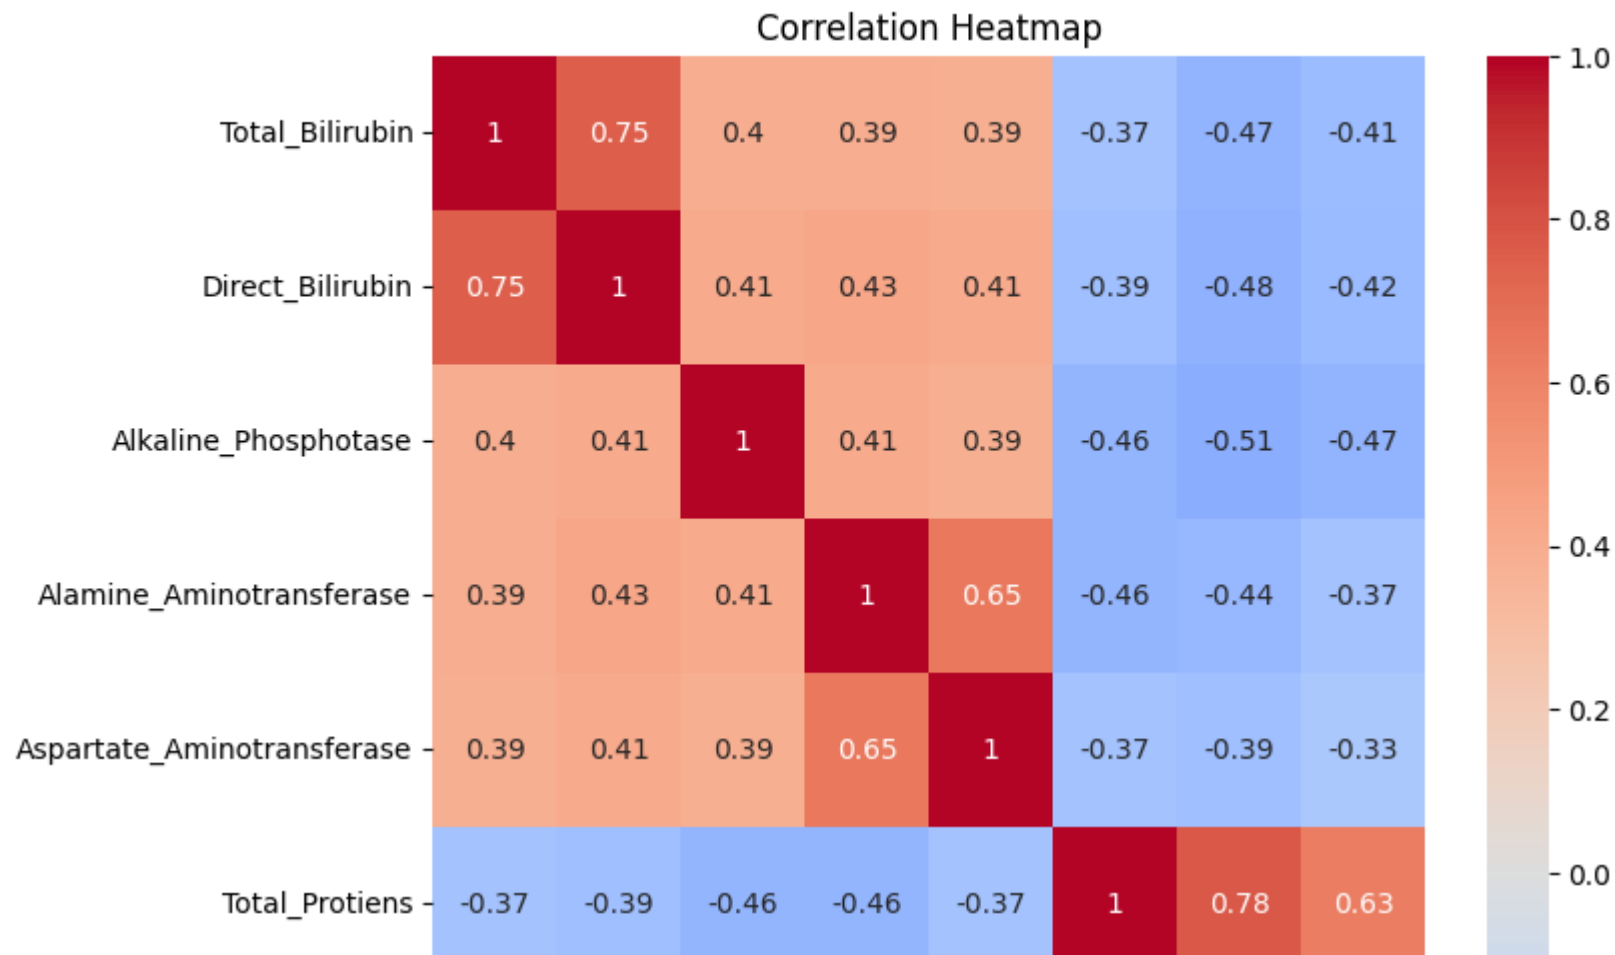
<Axes: >




✓ Finding outlier using heatmap

```
corr_matrix = X.corr()

plt.figure(figsize=(8, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', center=0)
plt.title('Correlation Heatmap')
plt.show()
```



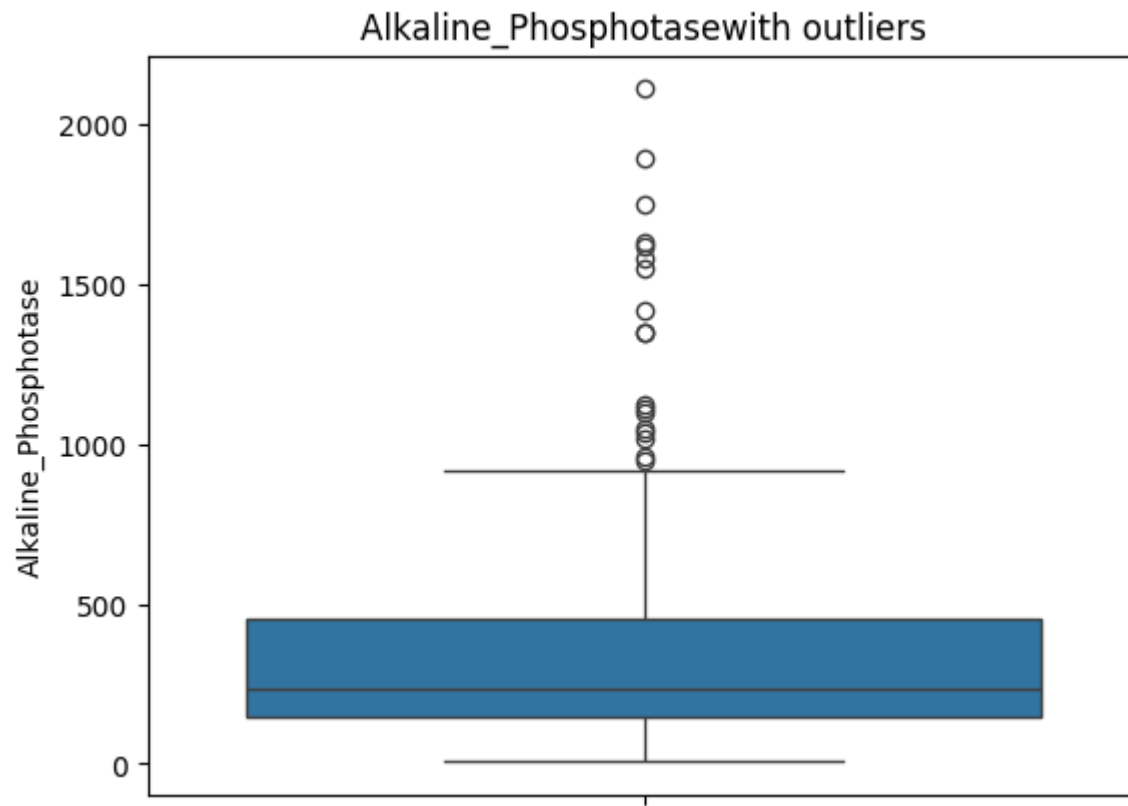


Double-click (or enter) to edit


```
# function for box plot
def printBox(df, col, title):
    sns.boxplot(df[col])
    plt.title(col + title)
```


✓ Finding outlier using box plot

```
printBox(df, 'Alkaline_Phosphatase', 'with outliers')
```



```
printBox(df, 'Aspartate_Aminotransferase', 'with outliers')
```





✓ Removing outlier using IQR

```
# Function to remove outliers
def removeOutlier(df, col):
    Q1 = df[col].quantile(0.25)
    Q3 = df[col].quantile(0.75)
    IQR = Q3 - Q1

    lower = Q1 - 1.5 * IQR
    upper = Q3 + 1.5 * IQR

    filtered_col = df[col][(df[col] >= lower) & (df[col] <= upper)]
    # filtered_col = (df[col] >= lower) & (df[col] <= upper)

    return filtered_col

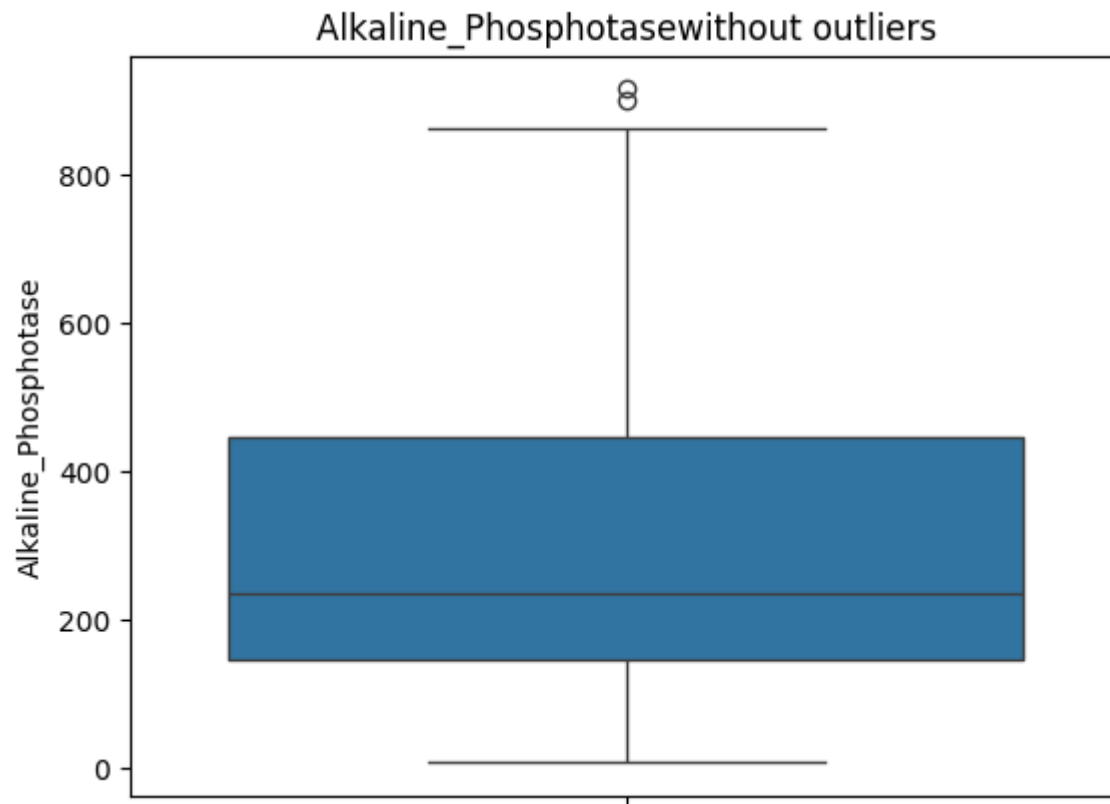
df['Alkaline_Phosphotase'] = removeOutlier(df, 'Alkaline_Phosphotase')

df['Aspartate_Aminotransferase'] = removeOutlier(df, 'Aspartate_Aminotransferase')
```

```
# df=df.copy()
```


✓ Data without outliers

```
printBox(df, 'Alkaline_Phosphotase', 'without outliers')
```



```
printBox(df, 'Aspartate_Aminotransferase', 'without outliers')
```

Aspartate_Aminotransferasewithout outliers



- ✓ Splitting dataset into training and testing data using train test split

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42, test_size = 0.2)
```


- ✓ Handle imbalanced data using SMOTE

```
# y_train.value_counts()[0], y_train.value_counts()[1]
```



```
# # ration = majority / minority  
# imbalance_ratio = y_train.value_counts()[1] / y_train.value_counts()[0]  
# imbalance_ratio
```

```
sns.countplot(x=y_train)  
plt.title('Class Distribution Before SMOTE')  
plt.show()
```



```
# !pip install imblearn
```

```
# from imblearn.over_sampling import SMOTE
# smote = SMOTE(random_state=42)

# X_train, y_train = smote.fit_resample(X_train, y_train)
```

```
y_train.value_counts()
```

	count
Dataset	
1	1347
0	565

dtype: int64


```
imbalance_ratio = y_train.value_counts()[1] / y_train.value_counts()[0]
imbalance_ratio
```


```
np.float64(2.384070796460177)
```

```
y_train.shape
```

```
(1912,)
```

```
sns.countplot(x=y_train)
plt.title('Class Distribution After SMOTE')
plt.show()
```





✓ Standardization using standard scaler

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
```

```
X_train_std = scaler.fit_transform(X_train)
```

```
X_test_std = scaler.transform(X_test)
```

✓ Evaluation Metrics, Loss and ROC curve functions

```
# Function to plot ROC curve
def rocCurve(y_test, y_pred):
    from sklearn.metrics import roc_curve, auc
    fpr, tpr, thresholds = roc_curve(y_test, y_pred)
    roc_auc = auc(fpr, tpr)

    plt.figure()
    plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' % roc_auc)
    plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.05])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('Receiver Operating Characteristic (ROC)')
    plt.legend(loc="lower right")
    plt.show()
```

```
# Function to plot loss
def plot_loss(training_loss, validation_loss):
    epochs = range(1, len(training_loss) + 1)
    plt.plot(epochs, training_loss, 'r', label='Training Loss')
    plt.plot(epochs, validation_loss, 'b', label='Validation Loss')
    plt.title('Training and validation Loss')
    plt.xlabel('Epochs')
    plt.ylabel('Loss')
    plt.legend()
    plt.show()
```

```
# Function for Evaluation Metrics
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
```

```

def calculate_metrics(model, X_train_std, X_test_std, y_train, y_test):

    y_train_pred = model.predict(X_train_std)
    y_train_pred_labels = (y_train_pred > 0.5).astype(int)

    training_accuracy = accuracy_score(y_train, y_train_pred_labels)
    training_precision = precision_score(y_train, y_train_pred_labels)
    training_recall = recall_score(y_train, y_train_pred_labels)
    training_f1 = f1_score(y_train, y_train_pred_labels)

    print(f"Training Accuracy: {training_accuracy}")
    print(f"Training Precision: {training_precision}")
    print(f"Training Recall: {training_recall}")
    print(f"Training F1 Score: {training_f1}")

    y_pred = model.predict(X_test_std)
    y_pred_labels = (y_pred > 0.5).astype(int)
    accuracy = accuracy_score(y_test, y_pred_labels)
    precision = precision_score(y_test, y_pred_labels)
    recall = recall_score(y_test, y_pred_labels)
    f1 = f1_score(y_test, y_pred_labels)

    print(f"\nAccuracy: {accuracy}")
    print(f"Precision: {precision}")
    print(f"Recall: {recall}")
    print(f"F1 Score: {f1}")

    return y_pred

```

✎ Machine Learning algorithms

✓ KNIN

```
from sklearn.neighbors import KNeighborsClassifier
```

```
knn = KNeighborsClassifier(n_neighbors = 5)
```

```
knn_history = knn.fit(X_train_std, y_train)
```


✓ Result


```
knn_y_pred = calculate_metrics(knn, X_train_std, X_test_std, y_train, y_test)
```

```
Training Accuracy: 0.9435146443514645  
Training Precision: 0.9538461538461539  
Training Recall: 0.9665924276169265  
Training F1 Score: 0.9601769911504425
```

```
Accuracy: 0.9123173277661796  
Precision: 0.926605504587156  
Recall: 0.9439252336448598  
F1 Score: 0.9351851851851852
```

```
rocCurve(y_test, knn_y_pred)
```





✓ SVM

```
from sklearn.svm import SVC  
  
svm = SVC(kernel='linear')  
svm.fit(X_train_std, y_train)
```

▼ SVC ⓘ ?
SVC(kernel='linear')

✓ Result


```
svm_y_pred = calculate_metrics(svm, X_train_std, X_test_std, y_train, y_test)
```

```
Training Accuracy: 0.9189330543933054  
Training Precision: 0.944112262705205
```

Training Precision: 0.944115205785595
Training Recall: 0.9406087602078693
Training F1 Score: 0.942357753811826

Accuracy: 0.9206680584551148
Precision: 0.9353846153846154
Recall: 0.9470404984423676
F1 Score: 0.9411764705882353

```
rocCurve(y_test, svm_y_pred)
```



▼ Random Forest

```
from sklearn.ensemble import RandomForestClassifier
```

```
rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)  
rf_classifier.fit(X_train_std, y_train)
```

```
▼      RandomForestClassifier      ⓘ ?  
RandomForestClassifier(random_state=42)
```


▼ Result


```
rf_y_pred = calculate_metrics(rf_classifier, X_train_std, X_test_std, y_train, y_test)
```

```
Training Accuracy: 1.0  
Training Precision: 1.0  
Training Recall: 1.0  
Training F1 Score: 1.0
```

```
Accuracy: 0.9248434237995825  
Precision: 0.9331306990881459  
Recall: 0.956386292834891  
F1 Score: 0.9446153846153846
```

```
rocCurve(y_test, rf_y_pred)
```





✓ Save RF classifier model

```
import joblib
```

```
joblib.dump(rf_classifier, 'liver_rf_model.pkl')
```

```
['liver_rf_model.pkl']
```

✓ Creating and Training ANN model using keras

```
EPOCH = 50
```

Epoch = 50

```
import tensorflow as tf
tf.random.set_seed(3)
from tensorflow import keras
```

```
tf.__version__

'2.18.0'
```

```
num_features = X_train_std.shape[1]
print(num_features)

8
```

```
ann_model = keras.Sequential([
    keras.layers.Dense(10, input_shape=(num_features, ), activation='relu'),
    keras.layers.Dense(20, activation='relu'),
    keras.layers.Dense(40, activation='relu'),
    keras.layers.Dense(80, activation='relu'),
    keras.layers.Dense(1, activation='sigmoid')
])
```

```
ann_model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
history = ann_model.fit(x=X_train_std, y=y_train, validation_data=(X_test_std, y_test), epochs=EPOCH)
```

```
ann_training_loss = history.history['loss']
ann_validation_loss = history.history['val_loss']
```

[Show hidden output](#)


▼ Result

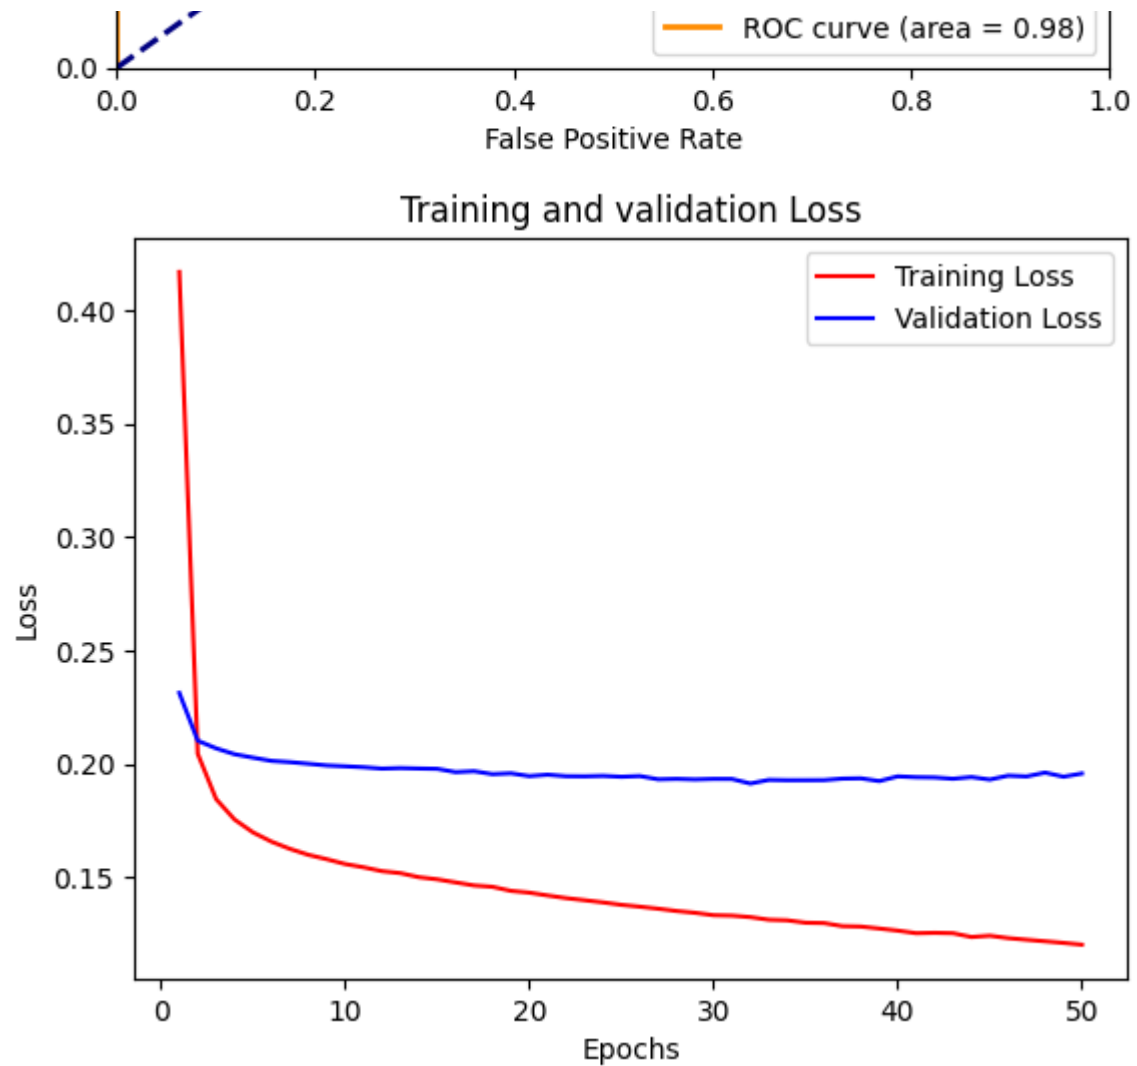
```
ann_y_pred = calculate_metrics(ann_model, X_train_std, X_test_std, y_train, y_test)
```

```
60/60 ————— 0s 4ms/step  
Training Accuracy: 0.9382845188284519  
Training Precision: 0.9658832448824868  
Training Recall: 0.9458054936896808  
Training F1 Score: 0.9557389347336834  
15/15 ————— 0s 11ms/step
```

```
Accuracy: 0.9018789144050104  
Precision: 0.928125  
Recall: 0.9252336448598131  
F1 Score: 0.9266770670826833
```

```
rocCurve(y_test, ann_y_pred)  
plot_loss(ann_training_loss, ann_validation_loss)
```





✓ Added l2 regularization

```
from tensorflow import keras
```

```
l2_model = keras.Sequential([
```

```
l2_model = keras.Sequential([
    keras.layers.Dense(10, input_shape=(num_features, ), activation='relu', kernel_regularizer=keras.regularizers.l2(0.02)),
    keras.layers.Dense(20, activation='relu', kernel_regularizer=keras.regularizers.l2(0.02)),
    keras.layers.Dense(40, activation='relu', kernel_regularizer=keras.regularizers.l2(0.02)),
    keras.layers.Dense(80, activation='relu', kernel_regularizer=keras.regularizers.l2(0.02)),
    keras.layers.Dense(1, activation='sigmoid')
])
```

```
l2_model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
history = l2_model.fit(x=X_train_std, y=y_train, validation_split=0.1, epochs=EPOCH)
```

```
l2_training_loss = history.history['loss']
l2_validation_loss = history.history['val_loss']
```

[Show hidden output](#)



```
l2_y_pred = calculate_metrics(l2_model, X_train_std, X_test_std, y_train, y_test)
```


```
60/60 ————— 0s 4ms/step
Training Accuracy: 0.922071129707113
Training Precision: 0.9450222882615156
Training Recall: 0.9443207126948775
Training F1 Score: 0.9446713702190865
15/15 ————— 0s 9ms/step
```

```
Accuracy: 0.9248434237995825
Precision: 0.9357798165137615
Recall: 0.9532710280373832
F1 Score: 0.9444444444444444
```

```
rocCurve(y_test, l2_y_pred)
plot_loss(l2_training_loss, l2_validation_loss)
```







✓ L1,L2 regularization

num_features

8

```
l1l2_model = keras.Sequential([
    keras.layers.Input(shape=(num_features,)),
    keras.layers.Dense(10, activation='relu'),
    keras.layers.Dense(20, activation='relu', kernel_regularizer=keras.regularizers.l1_l2(l1=0.01, l2=0.01)),
    keras.layers.Dense(40, activation='relu', kernel_regularizer=keras.regularizers.l1_l2(l1=0.01, l2=0.01)),
    keras.layers.Dense(80, activation='relu', kernel_regularizer=keras.regularizers.l1_l2(l1=0.01, l2=0.01)),
    keras.layers.Dense(1, activation='sigmoid')
])

l1l2_model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
history = l1l2_model.fit(x=X_train_std, y=y_train, validation_split=0.1, epochs=EPOCH)

l1l2_training_loss = history.history['loss']
l1l2_validation_loss = history.history['val_loss']
```

[Show hidden output](#)


```
l1l2_model.input_shape
```

```
(None, 8)
```

```
l1l2_y_pred = calculate_metrics(l1l2_model, X_train_std, X_test_std, y_train, y_test)
```

```
60/60 ————— 0s 4ms/step
```

```
Training Accuracy: 0.9173640167364017
```

```
Training Precision: 0.9426656738644825
```

```
Training Recall: 0.9398663697104677
```

```
Training F1 Score: 0.9412639405204462
```

```
15/15 ————— 0s 9ms/step
```

```
Accuracy: 0.9248434237995825
```


```
Precision: 0.9411764705882353
```



```
Recall: 0.9470404984423676
```

```
F1 Score: 0.9440993788819876
```

```
rocCurve(y_test, l1l2_y_pred)
```

```
plot_loss(l1l2_training_loss, l1l2_validation_loss)
```





✓ Added dropout layer

Add dropout layer

num_features

8

num_features = X_train_std.shape[1]

```
dropout_model = keras.Sequential([
    keras.layers.InputLayer(input_shape=(num_features,)),
    keras.layers.Dense(10, activation='relu'),
    keras.layers.Dropout(0.2),
    keras.layers.Dense(20, activation='relu'),
    keras.layers.Dropout(0.2),
    keras.layers.Dense(40, activation='relu'),
    keras.layers.Dropout(0.2),
    keras.layers.Dense(80, activation='relu'),
    keras.layers.Dropout(0.2),
    keras.layers.Dense(1, activation='sigmoid')
])
```

```
dropout_model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
history = dropout_model.fit(x=X_train_std, y=y_train, validation_split=0.1, epochs=EPOCH)
```

```
dropout_training_loss = history.history['loss']
dropout_validation_loss = history.history['val_loss']
```

[Show hidden output](#)


▼ Result

```
dropout_y_pred = calculate_metrics(dropout_model, X_train_std, X_test_std, y_train, y_test)
```

60/60 ————— **0s** 4ms/step
Training Accuracy: 0.9278242677824268
Training Precision: 0.9314775160599572
Training Recall: 0.9688195991091314
Training F1 Score: 0.9497816593886463
15/15 ————— **0s** 9ms/step

Accuracy: 0.9227557411273486
Precision: 0.9251497005988024
Recall: 0.9626168224299065
F1 Score: 0.9435114503816794

```
rocCurve(y_test, dropout_y_pred)  
plot_loss(dropout_training_loss, dropout_validation_loss)
```





✦ Using CNN

```
X_train_cnn = X_train_std.reshape(X_train_std.shape[0], X_train_std.shape[1], 1)
```

```
X_test_cnn = X_test_std.reshape(X_test_std.shape[0], X_test_std.shape[1], 1)
```

```
X_train_cnn.shape
```

```
X_train_cnn.shape
```

```
(1912, 8, 1)
```

```
cnnModel = keras.Sequential([
    keras.layers.Conv1D(filters=32, kernel_size=3, activation='relu', input_shape=(X_train_cnn.shape[1], 1)),
    keras.layers.MaxPooling1D(pool_size=2),

    keras.layers.Conv1D(filters=64, kernel_size=2, activation='relu', padding='same'),
    keras.layers.MaxPooling1D(pool_size=2),

    keras.layers.Flatten(),
    keras.layers.Dense(1, activation='sigmoid')
])

cnnModel.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

history = cnnModel.fit(X_train_cnn, y_train, epochs=EPOCH, validation_split=0.1, validation_data=(X_test_cnn, y_test))

cnn_training_loss = history.history['loss']
cnn_validation_loss = history.history['val_loss']
```

[Show hidden output](#)

```
cnn_y_pred = calculate_metrics(cnnModel, X_train_cnn, X_test_cnn, y_train, y_test)
```

```
60/60 ————— 1s 8ms/step
Training Accuracy: 0.9278242677824268
Training Precision: 0.9534883720930233
Training Recall: 0.9435783221974758
Training F1 Score: 0.9485074626865672
15/15 ————— 0s 15ms/step
```

```
Accuracy: 0.9206680584551148
Precision: 0.9300911854103343
Recall: 0.9532710280373832
F1 Score: 0.9415384615384615
```

```
rocCurve(y_test, cnn_y_pred)
plot_loss(cnn_training_loss, cnn_validation_loss)
```

[Show hidden output](#)

✓ Using multiple convolutional layer and pooling *layer*

```
cnnModel2 = keras.Sequential([
    keras.layers.Conv1D(filters=32, kernel_size=3, activation='relu', input_shape=(X_train_cnn.shape[1], 1), padding='
    keras.layers.MaxPooling1D(pool_size=2),

    keras.layers.Conv1D(filters=64, kernel_size=2, activation='relu', padding='same'),
    keras.layers.MaxPooling1D(pool_size=2),

    keras.layers.Conv1D(filters=128, kernel_size=2, activation='relu', padding='same'),
    keras.layers.MaxPooling1D(pool_size=1),

    keras.layers.Flatten(),
    keras.layers.Dense(1, activation='sigmoid')
])

cnnModel2.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

history = cnnModel2.fit(X_train_cnn, y_train, epochs=EPOCH, validation_split=0.1, validation_data=(X_test_cnn, y_test))

training_loss = history.history['loss']
validation_loss = history.history['val_loss']
```

Epoch 1/50

/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass
super().__init__(activity_regularizer=activity_regularizer, **kwargs)

60/60 **4s** 34ms/step - accuracy: 0.7834 - loss: 0.4668 - val accuracy: 0.9207 - val loss: 0.219

[illegible]


```

Epoch 22/50
60/60 ————— 0s 4ms/step - accuracy: 0.9217 - loss: 0.1431 - val_accuracy: 0.9186 - val_loss: 0.1822
Epoch 23/50
60/60 ————— 0s 4ms/step - accuracy: 0.9242 - loss: 0.1424 - val_accuracy: 0.9228 - val_loss: 0.1824
Epoch 24/50
60/60 ————— 0s 4ms/step - accuracy: 0.9226 - loss: 0.1418 - val_accuracy: 0.9228 - val_loss: 0.1822
Epoch 25/50
60/60 ————— 0s 4ms/step - accuracy: 0.9226 - loss: 0.1412 - val_accuracy: 0.9228 - val_loss: 0.1825
Epoch 26/50
60/60 ————— 0s 4ms/step - accuracy: 0.9228 - loss: 0.1408 - val_accuracy: 0.9228 - val_loss: 0.1834
Epoch 27/50
60/60 ————— 0s 4ms/step - accuracy: 0.9249 - loss: 0.1405 - val_accuracy: 0.9228 - val_loss: 0.1837
Epoch 28/50
60/60 ————— 0s 4ms/step - accuracy: 0.9249 - loss: 0.1401 - val_accuracy: 0.9228 - val_loss: 0.1838

```

```
cnn2_y_pred = calculate_metrics(cnnModel2, X_train_std, X_test_std, y_train, y_test)
```

```

60/60 ————— 1s 8ms/step
Training Accuracy: 0.9388075313807531
Training Precision: 0.9617117117117117
Training Recall: 0.9510022271714922
Training F1 Score: 0.9563269876819709
15/15 ————— 0s 17ms/step

```

```


Accuracy: 0.9311064718162839
Precision: 0.9390243902439024
Recall: 0.9595015576323987
F1 Score: 0.9491525423728814



```


```

rocCurve(y_test, cnn2_y_pred)
plot_loss(training_loss, validation_loss)

```







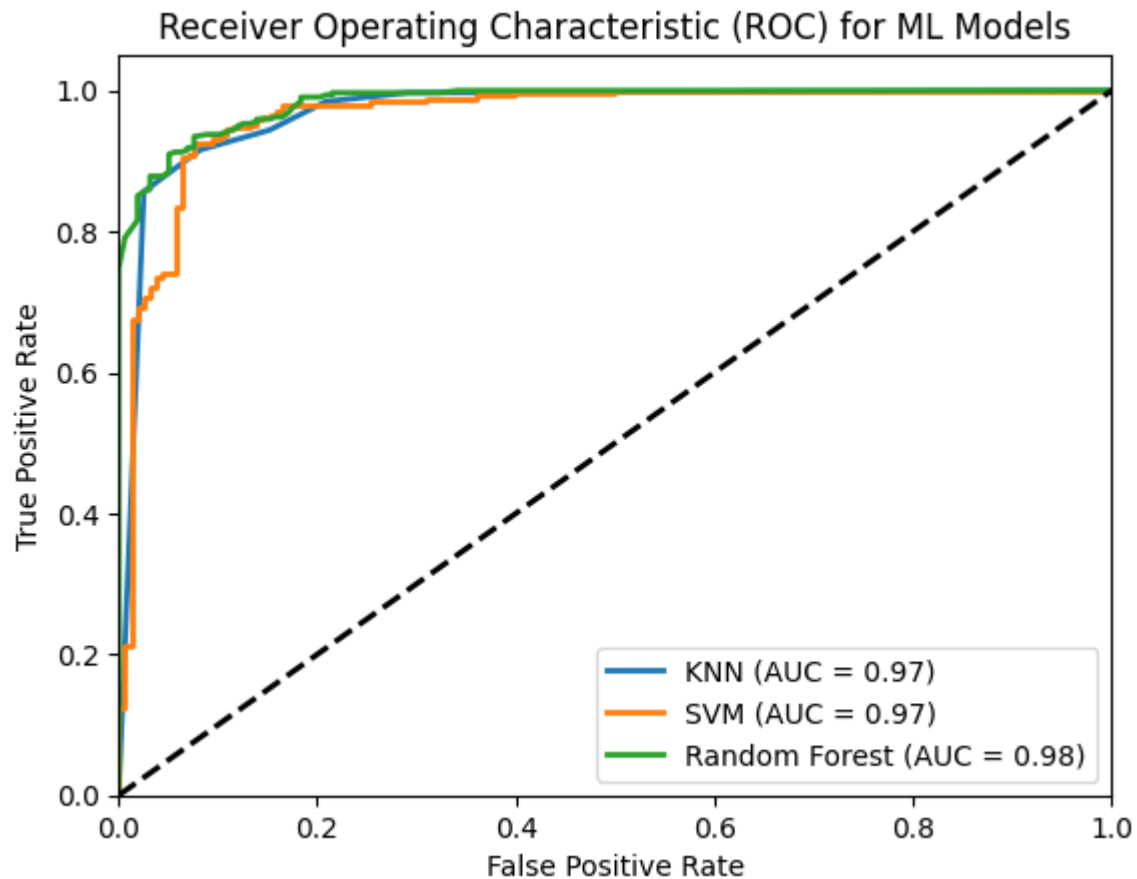
✓ Plot ROC curve for all the models

```
# plot ROC in single figure
from sklearn.metrics import roc_curve, auc
def plot_roc_curve(y_test, y_pred_prob, model_name):
    fpr, tpr, thresholds = roc_curve(y_test, y_pred_prob)
    roc_auc = auc(fpr, tpr)
    plt.plot(fpr, tpr, lw=2, label=f'{model_name} (AUC = {roc_auc:.2f})')
```

```
# ML
knn_y_pred_prob = knn.predict_proba(X_test_std)[: , 1]
svm_y_pred_prob = svm.decision_function(X_test_std)
rf_y_pred_prob = rf_classifier.predict_proba(X_test_std)[: , 1]

plot_roc_curve(y_test, knn_y_pred_prob, 'KNN')
plot_roc_curve(y_test, svm_y_pred_prob, 'SVM')
plot_roc_curve(y_test, rf_y_pred_prob, 'Random Forest')

plt.plot([0, 1], [0, 1], 'k--', lw=2)
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) for ML Models')
plt.legend(loc="lower right")
plt.show()
```




```
# DL
ann_y_pred_prob = ann_model.predict(X_test_std).ravel()
l2_y_pred_prob = l2_model.predict(X_test_std).ravel()
l1l2_y_pred_prob = l1l2_model.predict(X_test_std).ravel()
dropout_y_pred_prob = dropout_model.predict(X_test_std).ravel()
cnn_y_pred_prob = cnnModel.predict(X_test_cnn).ravel()
```


```
plot_roc_curve(y_test, ann_y_pred_prob, 'ANN')
plot_roc_curve(y_test, l2_y_pred_prob, 'ANN with L2')
plot_roc_curve(y_test, l1l2_y_pred_prob, 'ANN with L1L2')
plot_roc_curve(y_test, dropout_y_pred_prob, 'ANN with Dropout')
plot_roc_curve(y_test, cnn_y_pred_prob, 'CNN')
```

```
plot_roc_curve(y_test, cnn_y_pred_prob, 'CNN')
```

```
# plt.figure(figsize=(10, 8))
plt.plot([0, 1], [0, 1], 'k--', lw=2)
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) for DL Models')
plt.legend(loc="lower right")
plt.show()
```

15/15 ————— 0s 4ms/step
 15/15 ————— 0s 3ms/step
 15/15 ————— 0s 3ms/step
 15/15 ————— 0s 4ms/step
 15/15 ————— 0s 3ms/step

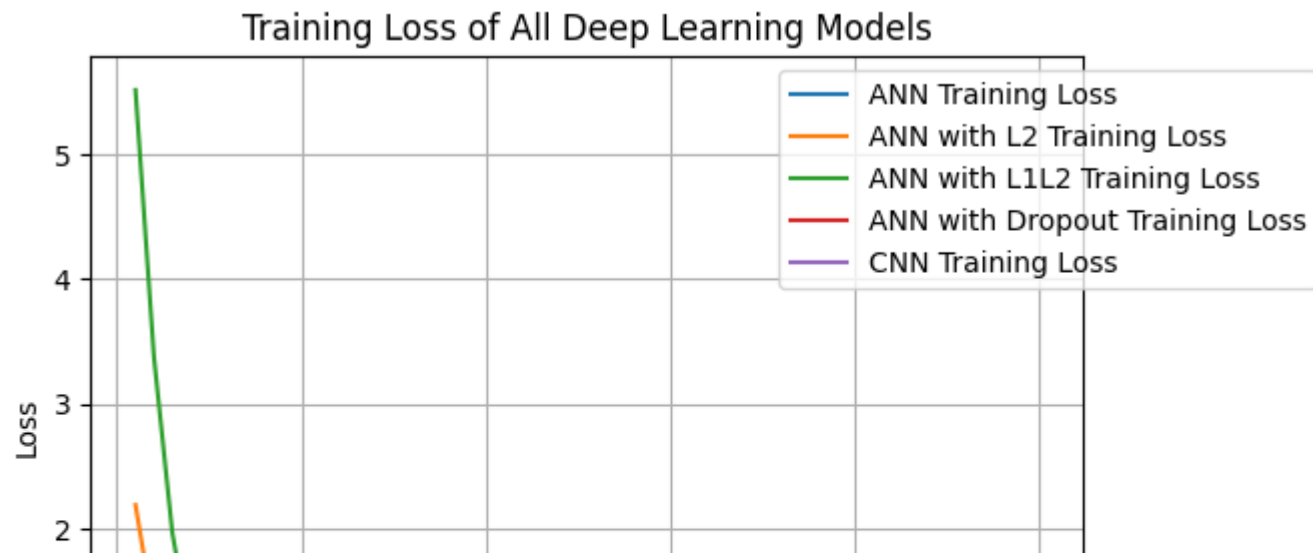





```
epochs = range(1, EPOCH + 1)
```

```
plt.plot(epochs, ann_training_loss, label='ANN Training Loss')
plt.plot(epochs, l2_training_loss, label='ANN with L2 Training Loss')
plt.plot(epochs, l1l2_training_loss, label='ANN with L1L2 Training Loss')
plt.plot(epochs, dropout_training_loss, label='ANN with Dropout Training Loss')
plt.plot(epochs, cnn_training_loss, label='CNN Training Loss')
```

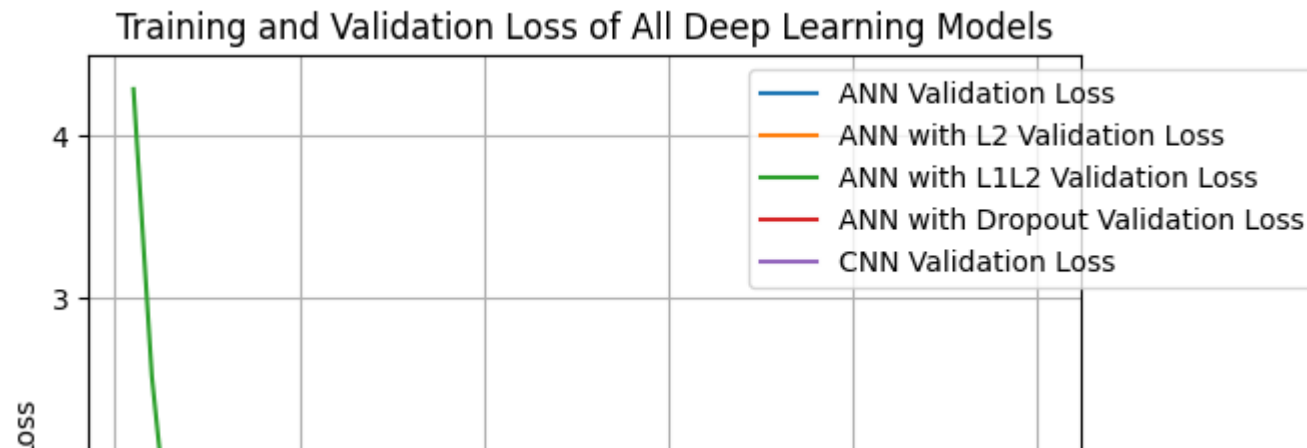
```
plt.title('Training Loss of All Deep Learning Models')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(loc='upper right', bbox_to_anchor=(1.25, 1))
plt.grid(True)
plt.show()
```






```
plt.plot(epochs, ann_validation_loss, label='ANN Validation Loss')
plt.plot(epochs, l2_validation_loss, label='ANN with L2 Validation Loss')
plt.plot(epochs, l1l2_validation_loss, label='ANN with L1L2 Validation Loss')
plt.plot(epochs, dropout_validation_loss, label='ANN with Dropout Validation Loss')
plt.plot(epochs, cnn_validation_loss, label='CNN Validation Loss')
```

```
plt.title('Training and Validation Loss of All Deep Learning Models')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(loc='upper right', bbox_to_anchor=(1.25, 1))
plt.grid(True)
plt.show()
```





✓ Plot Learning Curve for ML models

```
from sklearn.model_selection import learning_curve

def plot_learning_curve(estimator, title, X, y, ylim=None, cv=None,
                        n_jobs=None, train_sizes=np.linspace(.1, 1.0, 5)):

    plt.figure()
    plt.title(title)
    if ylim is not None:
        plt.ylim(*ylim)
    plt.xlabel("Training examples")
    plt.ylabel("Score")
    train_sizes, train_scores, test_scores = learning_curve(
        estimator, X, y, cv=cv,
        n_jobs=n_jobs,
        train_sizes=train_sizes)

    train_scores_mean = np.mean(train_scores, axis=1)
    train_scores_std = np.std(train_scores, axis=1)
    test_scores_mean = np.mean(test_scores, axis=1)
```



```

test_scores_std = np.std(test_scores, axis=1)
plt.grid()

plt.fill_between(train_sizes, train_scores_mean - train_scores_std,
                 train_scores_mean + train_scores_std, alpha=0.1,
                 color="r")
plt.fill_between(train_sizes, test_scores_mean - test_scores_std,
                 test_scores_mean + test_scores_std, alpha=0.1, color="g")
plt.plot(train_sizes, train_scores_mean, 'o-', color="r",
         label="Training score")
plt.plot(train_sizes, test_scores_mean, 'o-', color="g",
         label="Cross-validation score")


plt.legend(loc="best")
return plt


```

```


plot_learning_curve(knn, "Learning Curve (KNN)", X_train_std, y_train, cv=5)
plt.show()

```






```
plot_learning_curve(svm, "Learning Curve (SVM)", X_train_std, y_train, cv=5)  
plt.show()
```



```
plot_learning_curve(rf_classifier, "Learning Curve (Random Forest)", X_train_std, y_train, cv=5)
plt.show()
```



Start coding or [generate](#) with AI.

